

Azure Kubernetes Best Practices Workshop

Workbook Team 1

Preparation

clone the hands-on repository

```
git clone https://github.com/kubernetes-workshop/hands-on.git  
cd hands-on
```

please checkout your branch: team1

```
git checkout team1
```

Exercise #1

deploy to kubernetes

```
kubect1 run hello --image nginx --labels=app=hello --port 80 --namespace team1
kubect1 get pod -n team1
kubect1 port-forward pod/hello-xxxxxxxxx-xxxxx 8000:80
visit http://localhost:8000
```

create service with ClusterIP

```
kubect1 expose deployments hello --port 80 --type ClusterIP -n team1
kubect1 get services -n team1
kubect1 port-forward service/hello 8000:80
visit http://localhost:8000
```

set team1 as default namespace

```
kubect1 config set-context $(kubect1 config current-context) --namespace=team1
```

validate it

```
kubect1 config view
kubect1 get pod
kubect1 get pod -n team1
kubect1 get services
kubect1 get services -n team1
```

Exercise #2

set environment variables (use your Dockerhub username)

`$version="v1"`

`$username="your-dockerhub-username"`

build and publish to DockerHub

`cd podinfo`

`docker build . --tag $username/podinfo:$version`

`docker login --username $username --password xxxxxxxx`

`docker push $username/podinfo:$version`

get current manifests from 'hello' and create deployment.yaml and service.yaml

`kubectl get deployment hello -o yaml`

`kubectl get service hello -o yaml`

replace labels, selectors and images according to the new application:

`"name: podinfo"`

`"image: your-dockerhub-username/podinfo:v1"`

run in kubernetes

`kubectl apply -f deployment.yaml`

`kubectl apply -f service.yaml`

`kubectl get all`

port forward and visit <http://localhost:8000>

`kubectl port-forward service/podinfo 8000:80`

troubleshoot => ask for help

`kubectl describe pod/podinfo-xxxxxxxx-xxxx`

`kubectl logs pod/podinfo-xxxxxxxx-xxxx`

Exercise #3

create ingress, first take a look into the file and try to make sense of it

```
kubect1 apply -f ingress.yaml
```

give it a minute to install

```
visit http://team1.ddnss.de/
```

```
visit http://team1.ddnss.de/podinfo
```

play around (subdomain, wildcard, regex)

```
1) host: "podinfo.team1.ddnss.de"
```

```
2) path: /*
```

```
3) path: /foo/bar/[A-Z0-9]{3}
```

troubleshoot

```
kubect1 port-forward service/hello 8001:80
```

```
kubect1 port-forward service/podinfo 8002:80
```

```
visit http://localhost:8001
```

```
visit http://localhost:8002
```

Exercise #4

deploy fibo application

```
kubectl run fibo --image=fnbk/fibo --requests=cpu=200m --expose --port=80
```

```
kubectl autoscale deployment fibo --cpu-percent=50 --min=2 --max=10
```

```
kubectl get all
```

for each command open a new powershell (see scaling in action)

```
kubectl get hpa --watch
```

```
kubectl get pod --watch --selector run=fibo
```

loadtest manual

```
kubectl run --rm -it manual-loadtest --image=fnbk/loadtest /bin/bash
```

```
Curl http://fibo.team1.svc.cluster.local # check DNS resolution
```

```
/app/hey -z 3s -c 64 -m GET http://fibo.team1.svc.cluster.local # make 64 requests in 3 seconds
```

use a job to create an automated load test

```
kubectl apply -f loadtest.yaml
```

inspect, see what happens, see scaling in action

```
kubectl get all
```

```
kubectl describe job.batch/loadtest
```

```
kubectl logs pod/loadtest-job-xxxxx
```

cleanup loadtest job

```
kubectl get all
```

```
kubectl delete -f ./loadtest.yaml
```